

Gradient Boosting for Regression and Classification

--Summary and Understanding from StatQuest

	Regression	Classification	
	Solution, Action	Solution, Action	Understanding
Input data $\{(x_i, y_i)\}_{i=1}^n$, and a differentiable Loss Function $L(y_i, F(x_i))$	$F(x)$ is a weak model, Loss function uses scaled RRS $\frac{1}{2} \sum_i [y_i - F(x_i)]^2$	Loss function uses cross entropy $-\sum_i [p_i \ln(\hat{p}_i) + (1 - p_i) \ln(1 - \hat{p}_i)]$ log(odds): $\ln(odds) = -\ln\left(\frac{1-p}{p}\right)$ Probability: $p = \frac{1}{1 + e^{-\log(odds)}}$	Using those forms of loss functions are for later convenience; For classification, $F(x_i)$ is log(odds), log(odds) calculated by different models can add up together, but probabilities cannot.
Step1: the 0 th model is set to be a constant, with $F_0(x) = \underset{\gamma}{\operatorname{argmin}} \sum_i^n L(y_i, \gamma)$	Average over all y_i 's $\frac{\sum_i^n y_i}{n}$	Log(odds) of positive samples $\ln\left(\frac{\#ofpositives}{\#ofnegatives}\right)$	Minimize the loss function over all data points, using a constant estimate
Step2: for m=1 to M Where m denotes m th model	Iterate through weak models one by one, until a max number of models is reached or additional models don't improve the fit.		
Step2.1: compute pseudo residual $\gamma_{im} = -\left[\frac{\partial L(y_i, F(x_i))}{\partial F(x_i)}\right]_{F(x_i)=F_{m-1}(x_i)}$	$\gamma_{im} = y_i - F_{m-1}(x_i)$	$\gamma_{im} \approx p_i - \hat{p}_{m-1,i}$, when γ is small. $\hat{p}_{m-1,i}$ is the predicted probability that i th data point is positive, by model $F_{m-1}(x_i)$.	Calculate the pseudo residual for each data point, indicating where the previous model creates errors, preparing for building up the next model.

Step2.2: Fit a regression tree to γ_{im} values and denotes each leaf by R_{jm} , for all leaves in m^{th} tree.	Use mean squared error (across the leaves under the same node) as the metric to choose the best feature to split on.		Gather larger pseudo residuals together in a leaf, smaller pseudo residuals together in a leaf.
Step 2.3: For all j in m^{th} tree, compute $\gamma_{jm} =$ $\underset{\gamma}{\operatorname{argmin}} \sum_{x_i \in R_{jm}} L(y_i, F_{m-1}(x_i) + \gamma)$	Average of pseudo residuals that end up in the same leaf $\gamma_{jm} = \operatorname{ave}(\gamma_{im} \in R_{jm})$	$\gamma_{jm} = \frac{\sum(\gamma_{im} \in R_{jm})}{\sum_{x_i \in R_{jm}} [\hat{p}_{m-1,i} \times (1 - \hat{p}_{m-1,i})]}$ $\hat{p}_{m-1,i}$ is the predicted prob by model $F_{m-1}(x_i)$.	To approximate the values in the same leaf by a constant; iterate over all leaves.
Step 2.4: Update model $F_m(x) = F_{m-1}(x) + \eta \cdot \gamma_{jm} \cdot I(x \in R_{jm}) ,$ Where η is a scaling factor for each model, similar to step size	Add up all scaled results from a series of weak models.	Add up all scaled log(odds) from a series of weak models. Still need to convert to probability.	This is why it's called an ensemble method.
Step 3: Output $F_M(x)$	For new input x , use model $F_M(x)$ to predict.	Convert the result form the last step to probability by logistic function $p = \frac{1}{1 + e^{-\log(odds)}}$	May need to tune η to find the best converged result.